Advance Question – Senior Data Engineer

1. Advanced Architecture & Internals

- **What is Log Compaction?**
    - **Answer:** A mechanism where Kafka ensures that it always retains at least the last known value for each message key within the log of data for a single topic partition. It is useful for restoring state after a crash (e.g., in a database changelog).
- **Why is Kafka faster than traditional message queues?**
    - **Answer:** It uses **Sequential I/O** (writing to the end of a file is faster than random access), **Zero Copy** (moving data directly from disk to network buffer without copying to application memory), and heavy use of the **OS Page Cache**.
- **What is the role of the Controller Broker?**
    - **Answer:** It is a single broker in the cluster responsible for managing partition states, electing partition leaders, and monitoring for broker failures. If the controller fails, a new one is elected via ZooKeeper (or KRaft).

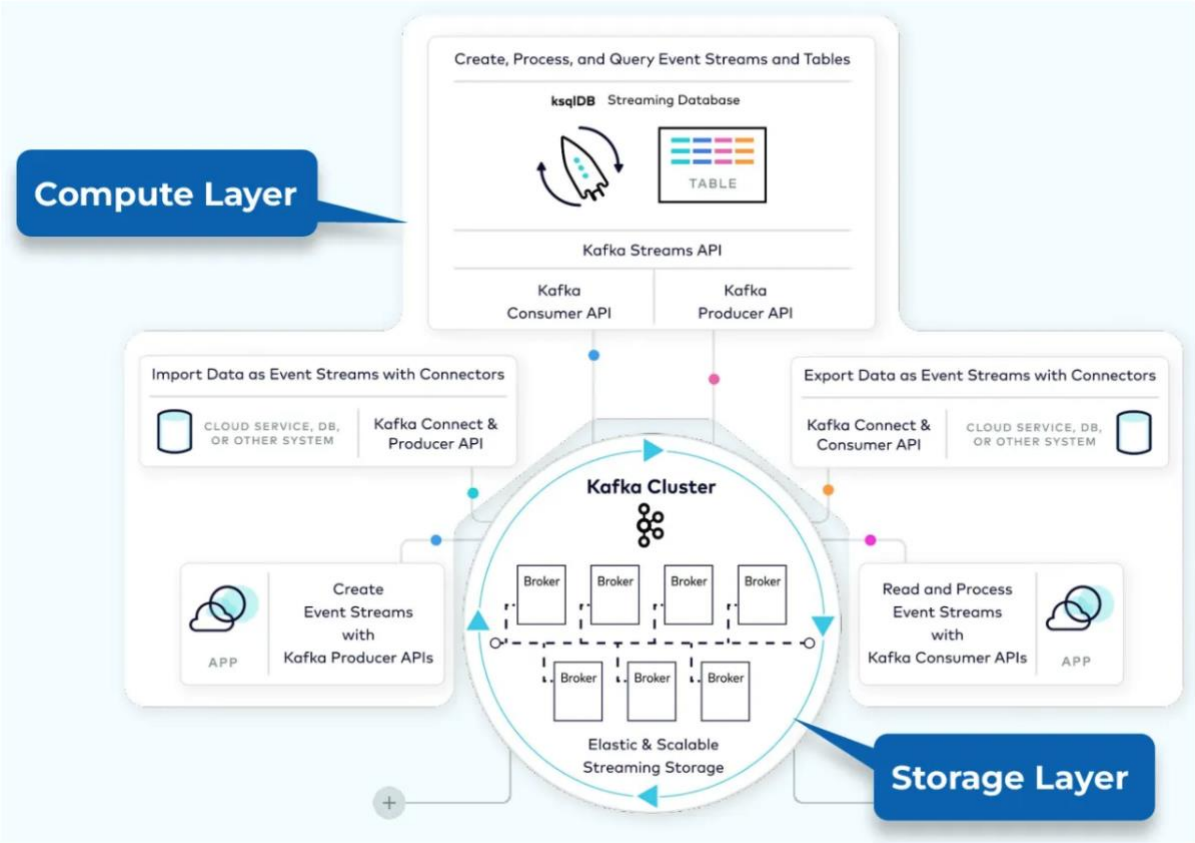

---

2. Reliability & Delivery Semantics

- **How do you achieve "Exactly-Once Semantics" (EOS)?**
  - **Answer:** This requires three configurations:
    1. **Idempotent Producer:** enable.idempotence=true (assigns sequence numbers to messages to prevent duplicates).
    2. **Transactional API:** Using initTransactions(), beginTransaction(), and commitTransaction().
    3. **Consumer Isolation:** Setting isolation.level=read_committed so consumers only see finalized transactions.
- **What is the "ISR shrinking" problem?**
  - **Answer:** It happens when "In-Sync Replicas" fall behind the leader.
  - **Causes:** Network latency, slow disk I/O on follower brokers, or the leader being overloaded with requests.

- **What is an "Unclean Leader Election"?**
    - **Answer:** When the leader of a partition fails and no "In-Sync" follower is available, Kafka can be configured to elect a "non-in-sync" follower. This ensures **availability** but risks **data loss**.

---

3. Scaling & Optimization Scenarios

- **Scenario: You have 10 partitions and 15 consumers in a single group. What happens?**
    - **Answer:** 10 consumers will each read from one partition, and **5 consumers will remain idle**. To fix this, you must increase the partition count to 15.
- **Scenario: Your consumer is experiencing high lag. How do you troubleshoot?**
    - **Answer:**
        1. Check if the number of partitions equals the number of consumers (Horizontal Scaling).
        2. Check for **Hot Partitions** (one partition has way more data due to a bad key).
        3. Increase fetch.min.bytes or max.poll.records to batch more data.
        4. Optimize the consumer's processing logic (e.g., moving DB writes to a background thread).
- **Scenario: You need to send a 50MB file through Kafka. Is this a good idea?**
    - **Answer:** No. Kafka is optimized for small messages (usually < 1MB).
    - **Solution:** Store the file in S3/Cloud Storage and only send the **Metadata/URL** through Kafka.

---

4. System Design & Integration

- **What is the Kafka Schema Registry?** *
    - **Answer:** A service that stores versioned schemas (usually Avro or Protobuf). It ensures that the producer and consumer are using compatible data

formats, preventing the "poison pill" message problem where a consumer crashes because it can't parse a message.

- **Kafka vs. RabbitMQ: When do you choose which?**
    - **Answer:** Choose **Kafka** for high-volume log processing, stream processing, and data replayability. Choose **RabbitMQ** for complex routing, priority queues, and simple request/response patterns.

---

5. Security & Operations

- **How do you secure data in Kafka?**
    - **Answer:** Use **SSL/TLS** for encryption in transit, **SASL** (Simple Authentication and Security Layer) for client-broker authentication, and **ACLs** (Access Control Lists) to restrict who can read/write to specific topics.

Scenario Based Questions :

1. The "Hot Partition" Problem

Scenario: You notice that one broker in your cluster has much higher CPU and disk usage than the others, and one specific partition has a massive lag while others are empty.

- **Cause:** This is usually due to **Data Skew**. The partition key being used (e.g., user_id) might have one specific value that appears way more often than others (a "Power User").
- **Solution:** * Change the partition key to something more granular.
    - Use a **Composite Key** (e.g., user_id + random_suffix).
    - If ordering isn't strictly required for that key, use a **Round Robin** strategy.

2. Consumer Group Rebalance Storm

Scenario: Your consumer group is constantly rebalancing, causing processing to stop every few minutes.

- **Cause:** The consumer is taking too long to process a single batch of messages. If the processing time exceeds max.poll.interval.ms, Kafka thinks the consumer is dead and triggers a rebalance.

- **Solution:** * Increase max.poll.interval.ms.
    - Decrease max.poll.records to process fewer messages per batch.
    - Optimize the processing logic (e.g., move database writes to an async thread).

3. Achieving "Exactly-Once" Semantics

Scenario: A financial application requires that no payment message is ever processed twice, even if the network fails during a retry.

- **Solution:** You must enable three things:
    1. **Producer:** Set enable.idempotence=true. This assigns a sequence number to messages so the broker can drop duplicates.
    2. **Transactions:** Use the Kafka Transactional API to ensure "all-or-nothing" writes.
    3. **Consumer:** Set isolation.level=read_committed so the consumer only reads successfully committed transactions.

4. Handling Large Messages (>1MB)

Scenario: A producer needs to send a 10MB image through Kafka, but the default limit is 1MB.

- **Solution:** * **The Claim-Check Pattern (Recommended):** Store the 10MB file in S3/Cloud Storage and send only the **Reference URL** through Kafka.
    - **In-Line (Not Recommended):** If you must, increase message.max.bytes on the broker and max.request.size on the producer, but be aware this can cause memory pressure.

5. Infinite Retries & Poison Pills

Scenario: A consumer keeps crashing on the same message because the data format is corrupted. It restarts, retries, and crashes again.

- **Solution:** Implement a **Dead Letter Queue (DLQ)**.
    - Catch the exception in the consumer code.
    - Send the "poison pill" message to a separate "error-topic" for manual inspection.

    - Commit the offset so the consumer can move to the next valid message.

6. Message Ordering with Multiple Partitions

Scenario: You need to process "Order Created" and "Order Paid" events in the correct sequence, but you have 10 partitions for scalability.

- **Solution:** Use the order_id as the **Partition Key**. Kafka guarantees ordering **within a single partition**. By using the same key, all events for a specific order are guaranteed to land in the same partition and be read in order.

7. Scaling Consumers Beyond Partitions

Scenario: You have a topic with 5 partitions and 5 consumers. Processing is still too slow. You add 5 more consumers to the group, but they sit idle.

- **Cause:** Kafka follows the rule: **One partition can only be read by one consumer in a group.** * **Solution:** You must **increase the number of partitions** to 10. Only then can the 5 idle consumers begin working.

8. Handling Broker Failure without Data Loss

Scenario: A broker crashes. When it comes back up, you realize some messages sent just before the crash are missing.

- **Cause:** The producer likely used acks=1 (only leader acknowledged) or acks=0.
- **Solution:** * Set acks=all so the producer waits for all In-Sync Replicas (ISR) to confirm.
    - Set min.insync.replicas=2 to ensure the message is written to at least two brokers before being considered "done."

9. High Consumer Lag During Traffic Spikes

Scenario: During a "Flash Sale," your consumer lag spikes to millions of messages.

- **Solution:** * **Temporary Scaling:** Increase partitions and consumers.
    - **Tuning:** Increase fetch.min.bytes and fetch.max.wait.ms to make fetching more efficient by pulling larger batches of data.

10. Transitioning from ZooKeeper to KRaft

Scenario: You are asked why modern Kafka (3.0+) is moving away from ZooKeeper.

- **Answer:** ZooKeeper creates a bottleneck for metadata (it can't handle 100k+ partitions easily). **KRaft** (Kafka Raft) stores metadata directly inside Kafka topics, allowing for much faster controller failover and support for millions of partitions.

Real interview Scenario Questions & Answers:

---

1. Message Duplication Scenario

Q: *Your consumer is processing the same message multiple times. How do you ensure idempotency?*

Answer:
Kafka gives **at-least-once delivery** by default, so duplicates can happen during retries, rebalances, or consumer restarts. I handle this by making the consumer **idempotent**.

My approach:

- Use a **unique business key / event ID** in each message.
- Before processing, check whether that event was already handled.
- Store processed IDs in a **dedup table**, cache, or target DB with a **primary key / unique constraint**.
- Make writes **upsert/merge-based** instead of blind inserts.
- If producing downstream Kafka messages, use **transactions** so consume-process-produce happens atomically.
- For compacted topics, I may use the latest key state to avoid repeated state updates.

So the key idea is: **duplicates may arrive, but business side effects should happen only once**.

---

2. High Lag in Consumer Group

Q: *Consumer lag suddenly spikes from 0 to 50k. What's your debugging approach?*

Answer:
I debug it layer by layer:

1. **Check if traffic increased**
    - Did producer throughput spike suddenly?
    - Is lag on all partitions or just a few?
2. **Look for partition skew**
    - One partition may be overloaded while others are idle.
3. **Check consumer health**
    - CPU, memory, GC pauses, thread starvation
    - Is the application stuck on one slow external dependency like DB/API?
4. **Check consumer settings**
    - max.poll.records
    - fetch.min.bytes
    - max.poll.interval.ms
    - commit behavior
5. **See if rebalances happened**
    - Frequent rebalances temporarily stop consumption.
6. **Check downstream bottlenecks**
    - Slow DB inserts, locking, network timeout, batch flush delays
7. **Broker-side validation**
    - Broker CPU, disk I/O, network, under-replicated partitions

My goal is to answer:
Is lag caused by more input, slower processing, or inability to keep partitions balanced?

---

3. Ordering Gone Wrong

Q: *Events for the same user come out of order. How do you ensure ordering?*

Answer:
Kafka guarantees ordering **only within a partition**. So if I need all events for the same user in order, I must ensure:

- The **partition key = user_id**
- All events for that user always go to the **same partition**

I also avoid breaking order by:

- Not using random partitioning for related events
- Avoiding unnecessary parallel processing inside the consumer
- Being careful with retries that can reorder downstream handling
- Limiting multiple in-flight producer requests when strict ordering matters

So the rule is:
correct partition key + single partition ordering + controlled consumer parallelism.

---

4. Producer Throughput Drops

Q: *Producer latency goes from 5ms to 200ms. What will you check?*

Answer:
I would check both producer tuning and broker health.

Producer side:

- batch.size too small
- linger.ms too high or too low
- compression enabled and whether CPU is becoming bottleneck
- acks=all causing wait due to slow replicas
- retries increasing latency
- too many requests in flight

Broker / infra side:

- broker disk I/O saturation
- network latency between app and brokers

- ISR shrink causing slower acknowledgments
- broker CPU or page cache pressure

I'd also compare:

- Is latency high for all topics or only one?
- Did traffic pattern change?
- Any broker leader change recently?

So I'd isolate whether the issue is **producer config, broker replication, network, or disk**.

---

5. Consumer Group Rebalance Issue

Q: *Consumer group keeps rebalancing every few minutes. How do you fix it?*

Answer:
Frequent rebalances usually mean consumers are appearing dead or unstable.

I check:

- session.timeout.ms
- heartbeat.interval.ms
- max.poll.interval.ms

Common reason:

- Processing takes too long, so consumer doesn't call poll() in time
- App experiences GC pauses or CPU starvation
- Instances are crashing / restarting
- Autoscaling causes frequent join/leave behavior

Fixes:

- Increase max.poll.interval.ms if processing is long
- Reduce per-poll processing time
- Move heavy processing to worker threads
- Tune heartbeat/session timeout properly

- Use **cooperative rebalancing** if possible
- Stabilize deployments and autoscaling

Main idea:
keep consumers alive, polling regularly, and avoid long blocking work in the poll loop.

---

6. Kafka Cluster Disk Full

Q: *Broker disk usage reaches 85%. What's your rescue plan?*

Answer:
This is an operational risk because once disks fill up, brokers become unstable.

Immediate rescue:

- Check which topics are consuming most storage
- Reduce retention temporarily where safe
- Delete unused topics if approved
- Throttle producers if necessary
- Move partitions away from overloaded brokers
- Expand broker disks / add brokers

Then stabilize:

- Rebalance partitions across brokers
- Review replication factor
- Review retention policy and compaction policy
- Monitor log segment growth rate
- Forecast capacity before reaching critical threshold

So I'd do:
protect cluster first, free space safely, then fix capacity planning.

---

7. Exactly-Once Delivery Scenario

Q: *A financial workflow needs exactly-once processing. How do you design it?*

Answer:
For a financial workflow, I'd use Kafka's **exactly-once semantics** carefully, but I'd also explain that true exactly-once must include the full business boundary.

Design:

- **Idempotent producer** enabled
- Use **transactions**
- If using Kafka Streams, enable **EOS**
- Consume → process → produce in a single transactional flow
- Persist output with idempotent keys or transactional DB writes where required

Important nuance:
Kafka can give exactly-once for **Kafka-to-Kafka pipelines**, but if an external DB or payment system is involved, I still need:

- idempotent writes
- dedup by transaction ID
- transactional outbox / inbox patterns if needed

So the interview-safe answer is:
Idempotent producer + Kafka transactions + idempotent downstream systems.

---

8. Hot Partition Problem

Q: *One partition receives 80% of the data. How do you solve it?*

Answer:
That means the partition key is skewed.

I would:

- Revisit the keying strategy
- Avoid low-cardinality keys
- Use a **composite key** like user_id + shard
- Increase partitions if overall throughput needs scaling

- If strict per-key ordering is needed, preserve ordering carefully while redistributing load

Important point:
Just increasing partitions does **not** fix skew if the same hot key still hashes to one partition.

So the real solution is:
better key design first, partitions second.

---

9. Slow Consumer Scenario

Q: *One slow consumer affects the whole pipeline. Why?*

Answer:
In a consumer group, each partition is owned by only one consumer at a time. If one consumer gets slow and owns important partitions:

- Lag builds on those partitions
- Offsets are committed late
- Downstream processing is delayed
- Overall pipeline SLA suffers even if other consumers are healthy

Reasons can include:

- slow DB/API dependency
- uneven partition assignment
- CPU/memory pressure
- bad batch size
- blocking logic

Fix:

- identify which partitions are lagging
- optimize processing
- reduce per-message overhead
- rebalance partition ownership

- increase partitions/consumers if scaling is needed

So one slow consumer matters because **partition ownership creates a bottleneck**.

---

10. Multi-Region Sync Scenario

Q: *Need real-time sync between two regions. What do you choose?*

Answer:
It depends on the architecture.

For Kafka-native replication across regions, I'd consider:

- **Cluster Linking**
- **MirrorMaker 2**

Then choose between:

- **Active-passive** for simpler disaster recovery
- **Active-active** only if conflict handling and failover design are very clear

Trade-offs:

- latency between regions
- topic failover strategy
- offset synchronization
- consumer locality
- duplicate/conflict risk in active-active setups

For most business systems, I'd prefer **active-passive** first because it is operationally safer. I'd choose active-active only when the application can clearly handle distributed write complexity.

---

11. DLQ Scenario

Q: *Some messages keep failing during consumption. How do you handle them?*

Answer:
I never let poison messages block the whole pipeline forever.

My pattern:

- Retry transient failures a few times
- Use **retry topics** with backoff
- Apply **exponential backoff**
- After retry limit, send the event to a **DLQ**
- Store error reason, timestamp, partition, offset, and payload metadata
- Monitor DLQ volume and create replay tooling

I also distinguish between:

- **Transient issues** → retry
- **Bad data / schema mismatch / permanent business rule failure** → DLQ faster

That keeps the main pipeline healthy while preserving failed events for later analysis.

---

12. Kafka Streams vs Consumer API

Q: *You need joins + windowing + state. Choose Streams or Consumers?*

Answer:
I would choose **Kafka Streams**.

Reason:

- Built-in **state stores**
- Native **windowing**
- Stream-stream and stream-table **joins**
- Easier handling of **event time**
- Better support for **exactly-once semantics**

I'd choose raw Consumer API only if:

- the logic is very custom
- I need complete manual control
- the use case is simple consumption without advanced stateful processing

So for joins, windowing, aggregations, and stateful workflows, **Streams is the better fit**.

---

13. Data Loss Scenario

Q: *Producer got ACK but consumer never saw data. Where can loss happen?*

Answer:
I'd inspect the delivery guarantee path carefully.

Possible reasons:

- Producer used acks=1, leader acknowledged but replicas did not persist
- ISR shrank and leader failed before followers caught up
- **unclean leader election** allowed data loss
- Message was written but removed later by **retention expiry**
- Consumer group skipped it due to offset commit issues
- Topic compaction removed older records for same key
- Wrong topic/environment was consumed

So "producer got ACK" does **not always mean durable across failures**, especially if durability settings are weak.

Best protection:

- acks=all
- proper replication factor
- min.insync.replicas
- disable unsafe leader election
- monitor retention and offsets

---

14. Kafka as Event Store

Q: *Team wants Kafka as an event store. What constraints will you highlight?*

Answer:
Kafka can act like an event store, but I'd clarify its limits.

What Kafka is good at:

- append-only event history
- replay
- decoupled consumers
- temporal event streams

Constraints:

- retention is time/size based unless carefully managed
- compaction keeps latest key state, not full relational history
- not designed for ad hoc querying like a database
- point lookups are inefficient
- long-term history often needs **tiered storage / archival**
- rebuilding application state may need **snapshotting**
- Kafka is a strong **event backbone**, but not a full replacement for a database. Use it with **snapshots, archival storage, and query systems** where needed.